PARTICIPANT HANDOUT

AI for New Residents

ESSENTIAL CONCEPTS, SAFE USE, AND CLINICAL JUDGMENT

WORKSHOP DATE
April 27, 2026

AUTHOR
Brian Pitts, MD, MEHP

AI TOOL OUTPUT
Information and suggestions

CLINICAL JUDGMENT
You are responsible for the decision

SAFE, INFORMED PATIENT CARE
Better decisions, better outcomes

VERIFY Check accuracy and relevance > **DECIDE** Apply clinical judgment > **DOCUMENT** Record key points and rationale > **ESCALATE** Involve others when uncertain or high risk

INSIDE THIS PACKET

1	Case 1 Worksheet	Read the discharge summary
2	Case 2 Worksheet	The alert vs. the bedside
3	Case 3 Worksheet	The patient consulted AI
4	Commitment Worksheet	One AI use, one fact to verify, one risk
5	AI Tools Compared	Strengths, limits, and access
6	The Six-Part Prompt	Role, task, context, constraints, format, tone

REMEMBER:
AI can assist learning and workflow, but clinical judgment remains yours.

Think critically.
Stay curious.
Care safely.

PREPARED FOR SMALL-GROUP DISCUSSION AND APPLIED PRACTICE

CASE 1 WORKSHEET

CASE 1: READ THE DISCHARGE SUMMARY

DISCHARGE SUMMARY

Scenario:
You are the intern covering the wards.
You receive a discharge summary generated with AI. It looks polished and complete.
Before you sign off, what should you check?

1 What looks plausible?

What parts of the summary seem reasonable or accurate?

- _____
- _____

2 What is wrong or missing?

What errors, gaps, or inconsistencies do you notice?

- _____
- _____

3 What could harm the patient?

What in this summary could lead to misunderstanding, wrong decisions, or patient harm?

- _____
- _____

4 What should you verify?

What qpecific facts, data, or plans do you need to check in the chart or with the team?

- _____
- _____

5 What would you tell your senior?

What questions or concerns would you raise before finalizing the discharge?

- _____
- _____

Remember: Fluent writing ≠ accurate medicine.
Verify → Decide → Document.

CASE 2 WORKSHEET

CASE 2: THE ALERT vs. THE BEDSIDE

Scenario:

The EHR fires a high-risk sepsis alert.
Your patient looks well and wants to go home.
What should you think, check, and do next?

1 What does the alert suggest?

What is the tool flagging, and why might it be concerned?

- _____
- _____

2 What does the patient show?

What does the history, exam, and bedside assessment tell you?

- _____
- _____

3 What data might be artifact or misleading?

What data inputs could be wrong, incomplete, or not appropriate for this patient?

- _____
- _____

4 What should you do now?

What actions are appropriate at this moment?

- _____
- _____

5 When should you escalate?

When and to whom should you escalate this situation?

- _____
- _____

Remember: AI alerts inform, but do not replace clinical judgment.
Verify → Decide → Document → Escalate

CASE 3 WORKSHEET

CASE 3: THE PATIENT CONSULTED AI

Scenario:

Your patient has been asking an AI chatbot health questions.
They want to know if its advice is right and what to do.
How will you respond?
What will you affirm, correct, and guide?

1 What did the patient ask AI?

What information or advice did the patient receive?

- _____
- _____

2 What did AI get right?

What parts of the advice are accurate or helpful?

- _____
- _____

3 What is misleading, incomplete, or unsafe?

What could be wrong, missing, or potentially harmful?

- _____
- _____

4 How will you respond?

How will you explain things clearly and support the patient?

- _____
- _____

5 What will you document?

What key points of this conversation should be documented?

- _____
- _____

Remember: Patients may use AI. Your role is to listen, clarify, correct misconceptions, and partner in safe, evidence-based care.

Verify → Decide → Document → Escalate

Take-Home Commitments

1. One AI use I will verify.

2. One risk I will watch for.

3. One moment I will call for help.

AI Tools Compared.

Tool	Best For	Strengths	Limitations	Medical Training Relevance	Data Currency	Citations & Transparency	Cost / Access	Bottom Line
ChatGPT	General reasoning, explanations, studying, writing assistance	Versatile, excellent at explaining complex topics, great study aid	May hallucinate, not medical-specific, answers can lack depth or nuance	High for learning concepts, test prep, case discussions—verify clinical details	Up to June 2024	Provides citations when available, not exhaustive	Free (limited) or ChatGPT Plus (subscription)	Great all-rounder for learning and productivity—always verify.
Claude	Deep reasoning, ethical analysis, nuanced discussion	Strong reasoning and writing, thoughtful and cautious responses	May be overly cautious, fewer integrations than competitors	High for clinical reasoning practice and reflective discussions	Up to May 2024	Cites sources when available, clearer refusals	Free (limited) or Claude Pro (subscription)	Excellent for reasoning-heavy tasks and careful analysis.
Gemini	Quick answers, multimodal (text, images, docs), Google integration	Fast, up-to-date web access, handles images and files well	Quality varies with prompt; less nuanced depth at times	Moderate to High useful for quick lookups—verify clinical accuracy	Up to May 2024 (web access extends further)	Links to web sources, varies in quality	Free (limited) or Gemini Advanced (subscription)	Great for quick research and multimodal tasks. Double-check.
MedGemma	Medical Q&A, clinical reasoning support, education	Built for healthcare, strong medical accuracy, safety focus	Less general knowledge outside medicine, smaller ecosystem	High—designed for medical education and clinical reasoning	Up to 2024 (varies by release)	References medical literature (where available)	Free (research preview via Hugging Face / Google)	Purpose-built for medicine—highly relevant and trustworthy.
UpToDate	Evidence-based clinical decisions at the point of care	Expert-authored, peer-reviewed, comprehensive and trusted	Subscription required, not for casual browsing or brainstorming	Very High—gold standard for clinical information	Continuously updated	Extensive references with grading of evidence	Subscription (institutional or individual)	The clinician's first stop for trusted, up-to-date answers.
DynaMed	Clinical answers at the bedside, differential diagnosis	Concise, practical, great summaries, evidence-based	Subscription required, less depth than UpToDate	High—excellent for quick clinical support and teaching	Continuously updated	References provided with evidence levels	Subscription (institutional or individual)	Fast, reliable clinical answers when time matters.
Abridge	Clinical documentation, notes, visit summaries.	Exceptional speech recognition, saves time, integrates with EHRs	Not a clinical knowledge tool; accuracy depends on input	High for workflow and learning from well-structured notes	Continuously improved	Not a knowledge tool—no clinical citations	Subscription (clinician facing)	Powerful scribe AI that improves productivity and accuracy.
Rocky	Medical education, cases, flashcards, explanations	Built for learners, great UX, spaced repetition, cases	Not for real-time clinical decisions; content scope still growing	High—tailored for medical students and residents	Regularly updated	References provided for key content	Freemium or Pro (subscription)	Fantastic study companion for medicine learners.

Note: Always verify critical clinical information with trusted sources before applying to patient care.

THE SIX-PART PROMPT

A simple framework for clearer, better AI responses

1 ROLE

Define the AI's role or expertise.

Set the perspective.

What it does:
Sets the context and tone of the response.

Example:
"You are an experienced healthcare consultant."

2 TASK

Clearly state what you want the AI to do.

Be specific.

What it does:
Defines the action and focus.

Example:
"Analyze the following clinical scenario and provide insights."

3 CONTEXT

Provide relevant background information.

Add the details that matter.

What it does:
Gives the necessary background to understand the task.

Example:
"A 65-year-old patient with diabetes presents with fatigue and edema..."

4 CONSTRAINTS

Specify any limits, requirements, or boundaries.

Guide the guardrails.

What it does:
Helps the AI stay within boundaries.

Example:
"Use evidence-based sources from the last 5 years. Keep it under 300 words."

5 FORMAT

Define the structure or output format you want.

Shape the output.

What it does:
Ensures the response is easy to read and use.

Example:
"Provide the answer in bullet points with clear headings."

6 TONE

Set the tone or style of the response.

Match the audience.

What it does:
Aligns the language and style with your needs.

Example:
"Use a professional and supportive tone suitable for clinicians."

Put it all together:
When you include all six parts, you give the AI everything it needs to deliver accurate, relevant, and useful results.